



The Fabric Professional's AI Playbook: Skills, Models, and Survival Strategies

Dallas Fabric Tour
Sep 6, 2025

K. Scott Ferguson, PhD
Ferguson Analytics

The Reality Check

Your Role is Changing Whether You Adapt or Not

AI Accessibility Revolution

AI features now accessible at F2+ capacity (down from F64+)

Traditional Tasks Automated


Traditional analyst tasks increasingly automated

New Role Emerging

AI-Agent Orchestrator

The question isn't whether AI will change our work, but whether we'll lead that change or react to it. You're becoming an AI trainer and agent architect, not just a report builder.





Three Keys to AI Success

The Mindset Shift Every Professional Needs



Practice

Learning AI is Unlike Any Other Skill You've Learned. AI skills require continuous experimentation and iteration. It's not like learning Excel where you master functions once - AI responses vary and improve with practice.



The AI - Human Spectrum

Every Process Falls Somewhere Between All AI and All Human. Most valuable work happens in the hybrid zone where humans and AI collaborate. This includes tasks like data analysis, content creation, and strategic planning, where different tasks might fall on this spectrum.



Human Judgment

Embrace Human Judgment While Mastering AI Direction. We need to get really good at directing AI to do what we want. Human judgment becomes more valuable, not less, in an AI world. The skill is knowing when to trust AI and when to override it.

First Things First - Setup & Governance

Turn the Right Switches On

01

Enable Copilot

Enable Copilot for your tenant

02

Certify Models

Endorse/certify key semantic models

03

Mark Datasets

Mark datasets as **AI-Prepped**

04

Create Answers

Create Verified Answers for FAQs

05

Monitor Usage

Monitor usage and logs from day one

 If these aren't set, your AI will give poor answers and trust will crater. Governance enables trust, which enables adoption.



Presentation Agenda

1) Detailed Roadmap Illustration with Interccnnected Nodes and Pathways

What We'll Cover



Essential Prompting Skills

3 patterns that save hours



Model Transformation Demo

Make your data agent-ready



Your Career Action Plan

Positioning for the AI era

Everything we cover has detailed examples and templates in the GitHub playbook - both technical skills and career positioning strategies you can implement immediately.



Prompt Engineering - Revenue Analysis

✗ Bad Prompt

"Show me revenue trends"

Result: Vague chart, no context, no actionable insights

✓ Good Prompt

"Analyze Total Revenue (USD) for the last 6 months, identify the top 3 contributing factors to any significant changes, and suggest 2 specific actions we should consider based on the data. Include assumptions."

Result: Detailed breakdown, drivers identified, actionable recommendations

Context

Add timeframes

Verification

Ask for assumptions



Schema

Name specific measures

Action

Request recommendations



Prompt Engineering - Segmentation

The Business Intelligence Prompt Pattern

"Segment trips by payment type and fare totals. For each segment, provide the size, average trip value, and recommended strategy."

Specific Task

Clear segmentation request

1

Actionable Recommendations

Strategic insights, not just data

3

Multiple Dimensions

Payment type + fare totals

✓ This is how you add strategic value - you're generating insights, not just pulling raw data. This approach positions you as a business strategist, not just a data puller.

Model Transformation - Before

Traditional Semantic Model (Human-Designed)

Current Structure


- **Tables:** Trip, Medallion, HackneyLicense
- **Measures:** Revenue, AvgFare
- **Columns:** MedallionID, HackneyLicenseID, TotalAmount

Agent Query Example

"What's our revenue growth?"

Poor Response

"I found revenue data but I'm not sure which measure to use or what time period constitutes 'this quarter' in your business context."



This should feel familiar: designed for humans who already know the schema. Same data, but completely different AI experience coming up.





Model Transformation - After

AI-Ready Semantic Model

Transformed Structure

- **Tables:** Taxi Trips (Completed)
- **Measures:** Total Revenue (USD) – excludes no-charge and invalid trips
- **Measures:** Average Fare (USD)
- **Columns:** Trip Distance (miles), Trip Duration (sec), Payment Type

Same Query

"What's our revenue growth?"

Excellent Response

"Total Revenue (USD) for Q4 2025 is \$2.3M, representing 15% growth compared to Q3 2025 (\$2.0M) and 23% growth year-over-year (Q4 2024: \$1.87M). This excludes no-charge trips and invalid records."

- ✓ This is your new core skill - designing data that AI can understand. The dramatic improvement in response quality comes from descriptive names, business context, and clear measure definitions.

Key Design Principles

Making Data Agent-Friendly



Descriptive Names

Clear, business-friendly terminology that anyone can understand



Business Context

Capture inclusions/exclusions and business rules



Rich Descriptions & Synonyms

Feed Copilot context with comprehensive metadata



Consistent Conventions

Standardize naming and structure across all models



Verified Answers

Lock in consistent responses to frequently asked questions



Career Insight: You're becoming a **Data-AI Translator**. This translation skill between business language and AI understanding is your competitive advantage.

Common Pitfalls to Avoid

Anti-Patterns That Kill AI Success

Duplicate/Confusing Measures

Multiple revenue measures → AI chooses wrong one → executives lose trust

No Descriptions

AI guesses randomly without context → inconsistent answers


Too Many Sources

Information overload → confusion and poor performance

No Feedback Loop

Quality stagnates without continuous improvement

- ⊗ We had duplicate revenue measures in one implementation – Copilot chose the wrong one and executives lost trust immediately. Avoid these pitfalls on day 1 to maintain credibility.



Fabric Data Agents

01

Scope Carefully

Limit to ≤ 5 sources (ideally 1 semantic model)

02

Add Instructions

Include agent instructions for routing rules

03

Seed with Examples

Provide example Q&A pairs for training

04

Ensure Security

User identity passthrough, governed, read-only access

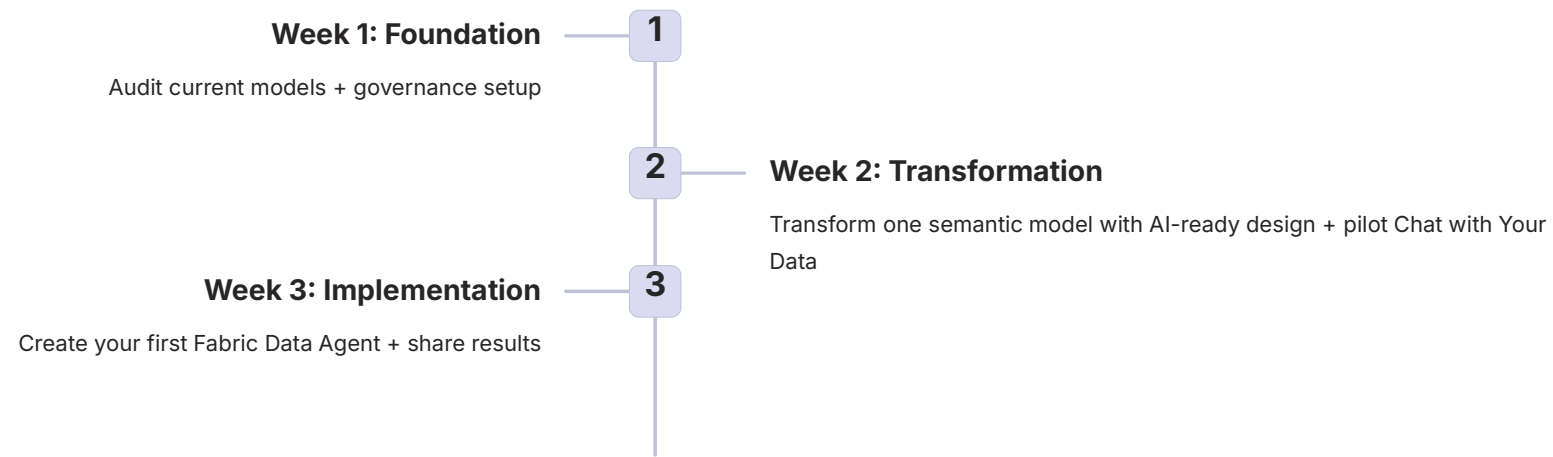
05

Deploy Everywhere

Surface in Power BI chat, Teams, Copilot Studio

Publishing agents to Teams and Copilot Studio creates a new distribution channel for your analytics - reaching users where they already work.

Your Career Evolution Action Plan



Career Positioning Strategy

LinkedIn Profile
Add "AI-Ready Data Architecture" to skills

Resume Update
"Designed semantic models for conversational analytics"

Project Leadership
Lead AI agent pilot initiatives

✔ These skills are in high demand - position yourself as an early adopter. Pick your most-used report, not your most complex one, to start.

Questions & Next Steps

Resources

- Playbook: <https://github.com/ksferguson/fabric-pro-ai-playbook>
- Communities: Local Power BI & Fabric User Groups
- Connect: <https://www.linkedin.com/in/kscottferguson>

Just get started!

